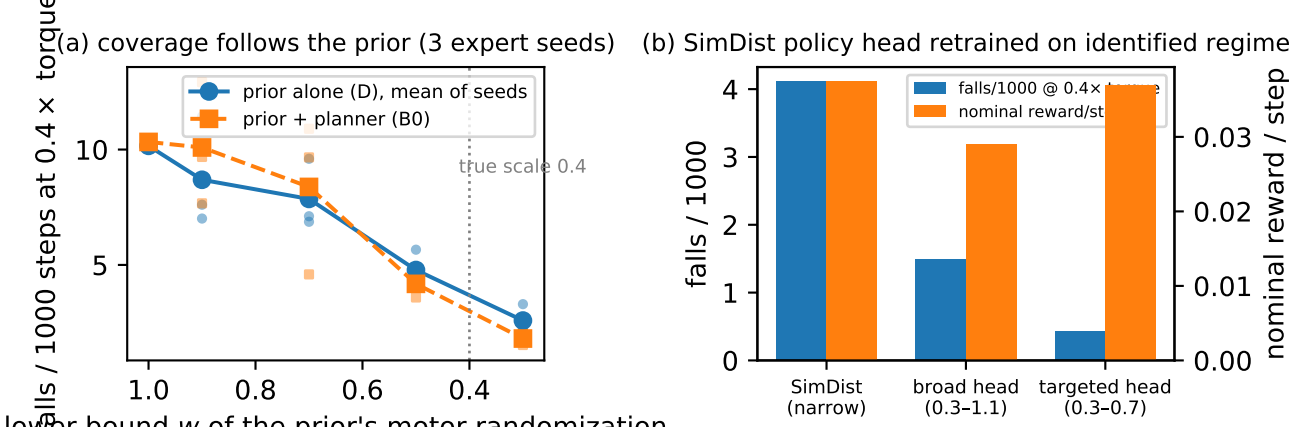
(a) coverage follows the prior (3 expert seeds)
falls / 1000 steps at 0.4 × torque
prior alone (D), mean of seeds
prior + planner (B0)
true scale 0.4
10
5
1.0
0.8
0.6
0.4
lower bound w of the prior's motor randomization
(b) SimDist policy head retrained on identified regime
falls / 1000
nominal reward / step
falls/1000 @ 0.4×
nominal reward/st
4
3
2
1
0
0.03
0.02
0.01
0.00
SimDist
(narrow)
broad head
(0.3–1.1)
targeted head
(0.3–0.7)